

Fork & zombies

Dr. Vimal Baghel
Assistant Professor
SCSET, BU

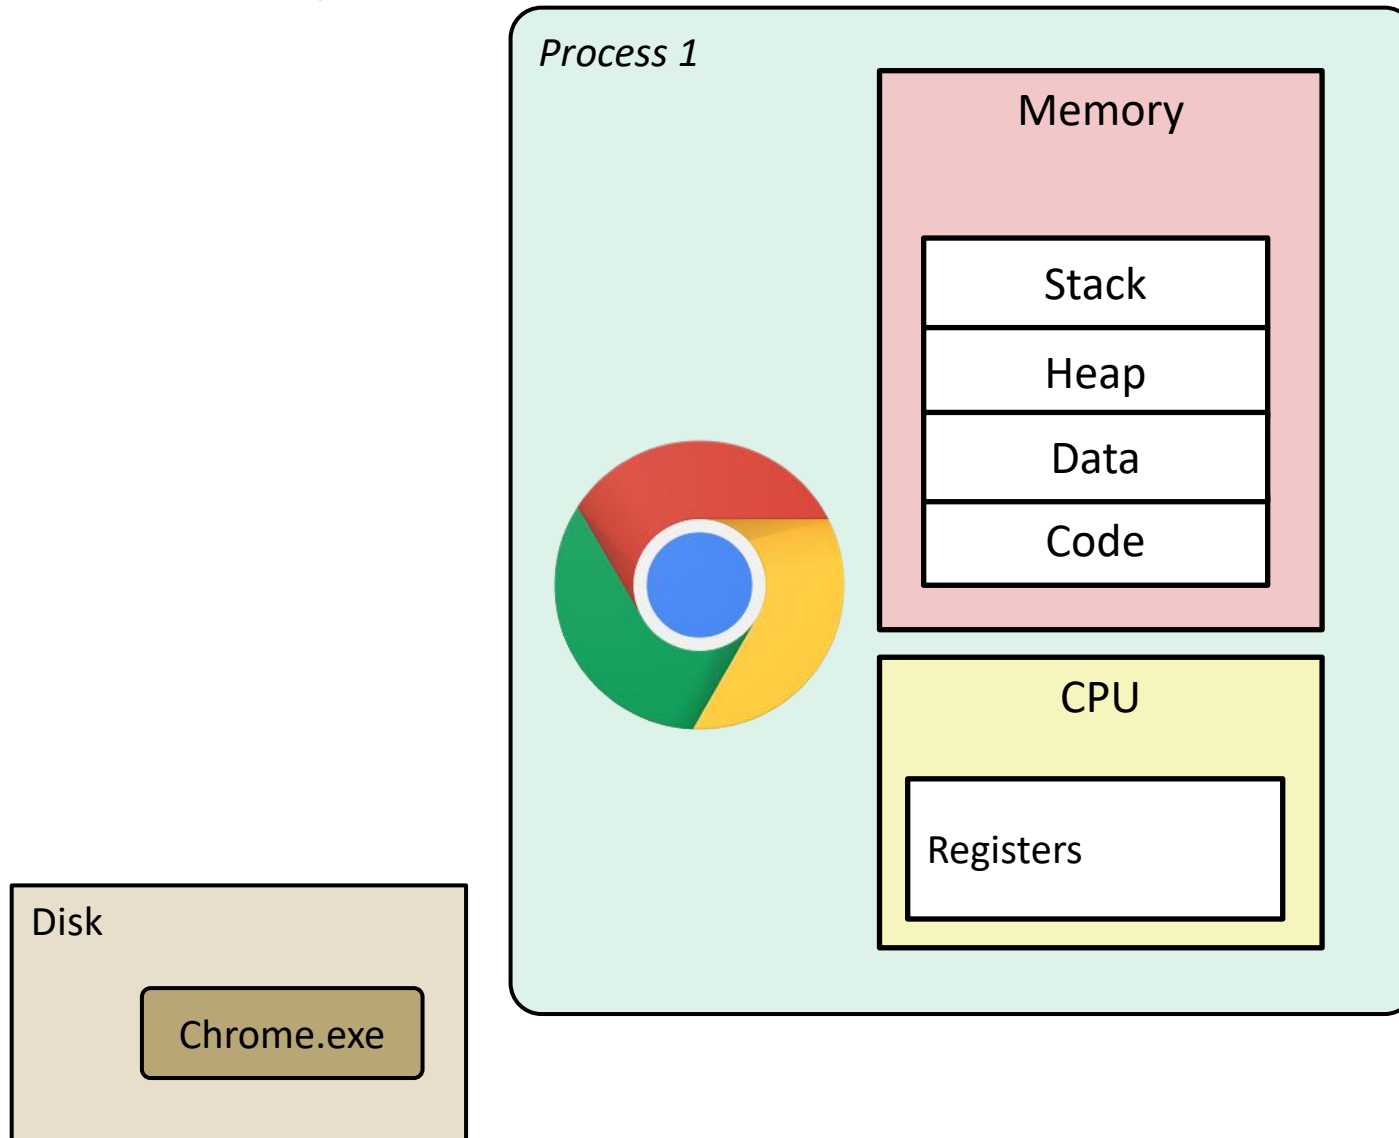


Processes

- **Processes and context switching**
- Creating new processes
 - `fork()`, `exec*()`, and `wait()`
- Zombies

What is a process?

It's an *illusion*!




What is a process?

- Another *abstraction* in our computer system
 - Provided by the OS
 - OS uses a data structure to represent each process
 - Maintains the *interface* between the program and the underlying hardware (CPU + memory)
- What do *processes* have to do with *exceptional control flow*?
 - Exceptional control flow is the *mechanism* the OS uses to enable **multiple processes** to run on the same system
- What is the difference between:
 - A processor? A program? A process?

hardware the "blueprint" an instance

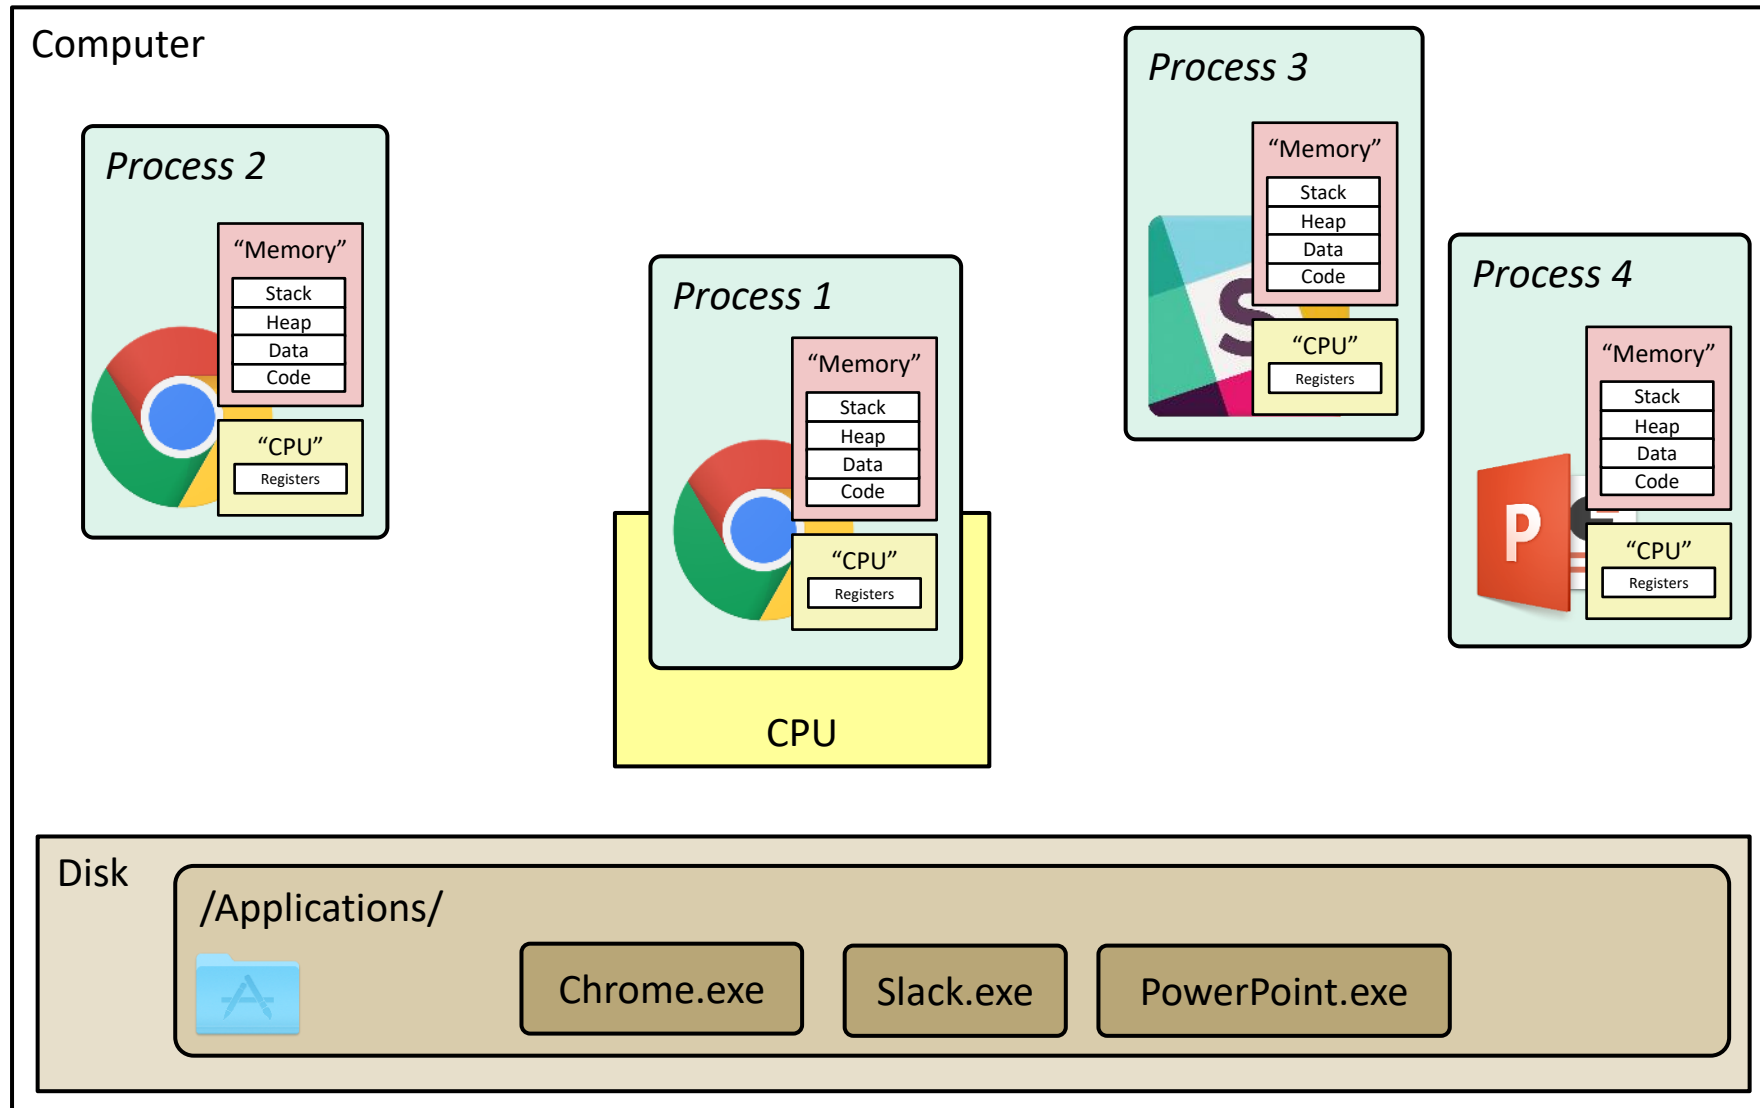
Processes

- A **process** is an instance of a running program
 - One of the most profound ideas in computer science
 - Not the same as “program” or “processor”
- Process provides each program with two key abstractions:
 - *Logical control flow*
 - Each program seems to have exclusive use of the CPU
 - Provided by kernel mechanism called **context switching**
 - *Private address space*
 - Each program seems to have exclusive use of main memory
 - Provided by kernel mechanism called **virtual memory**



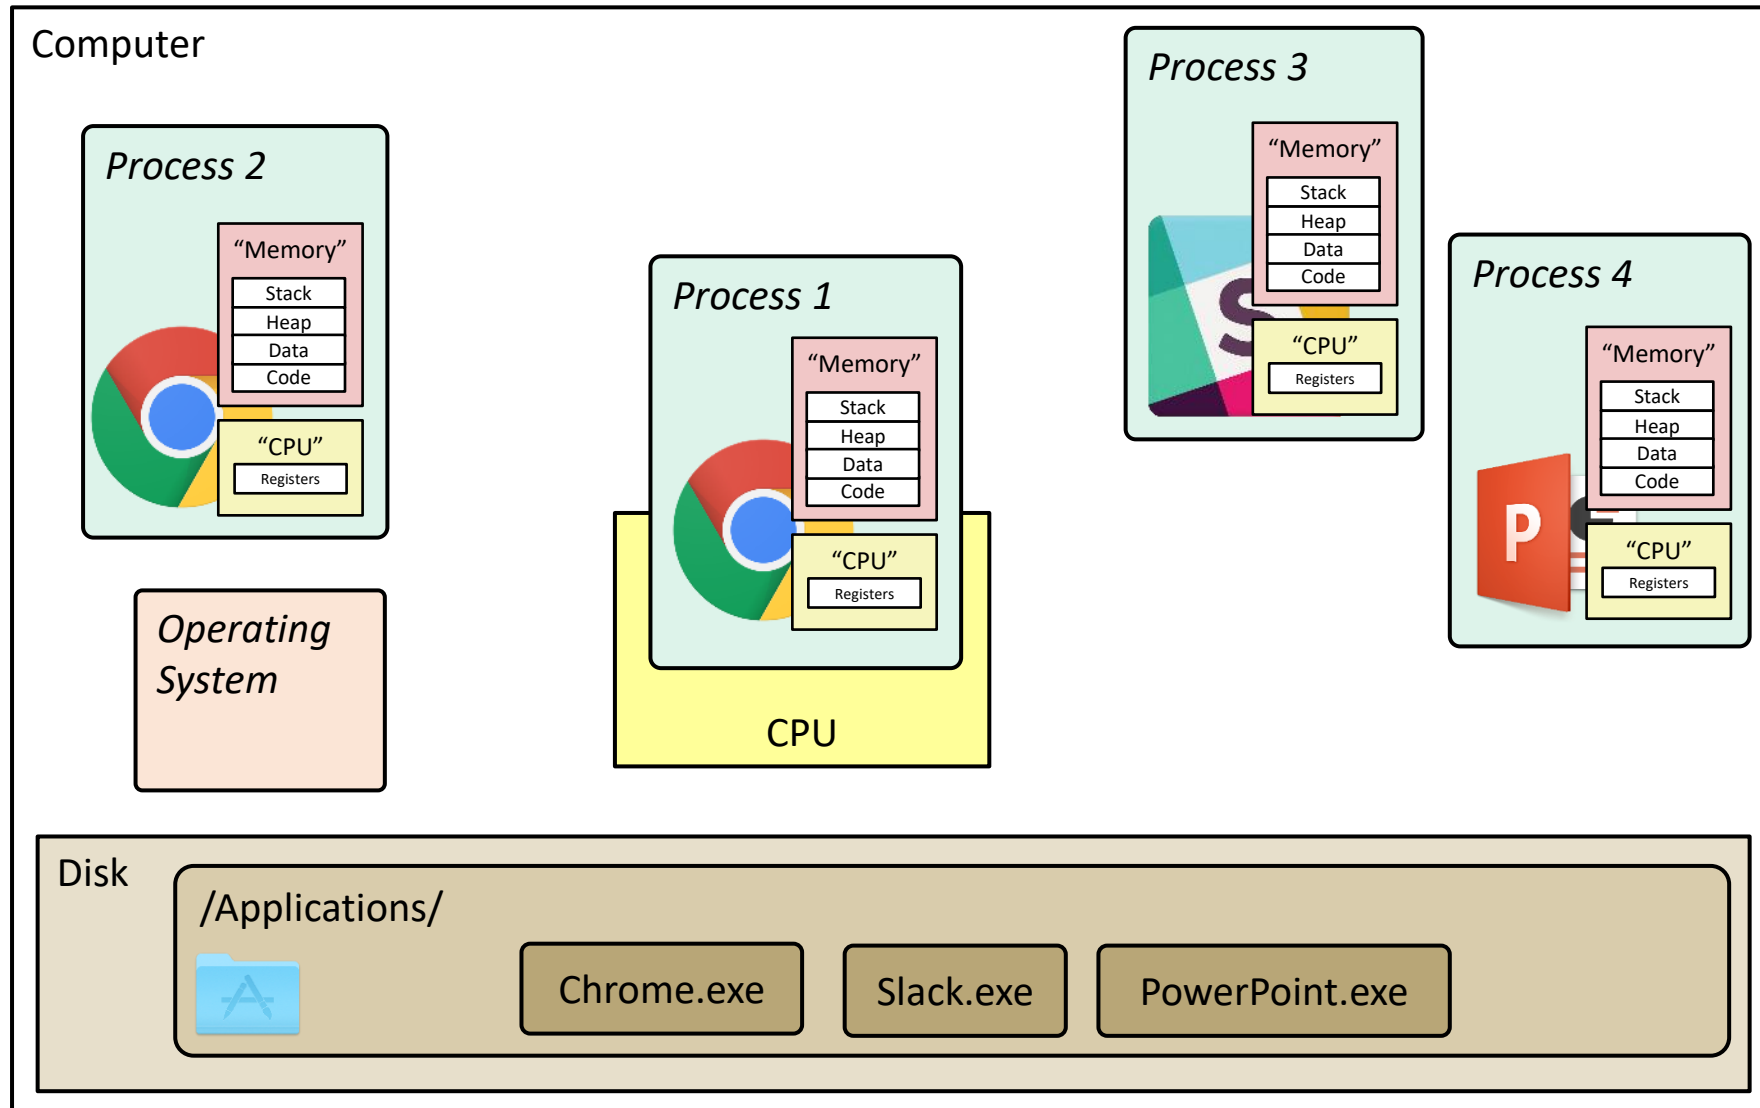
What is a process?

It's an illusion!



What is a process?

It's an *illusion*!



Multiprocessing on Uniprocessor: The Illusion


- While true multiprocessing is not possible on a uniprocessor, there are techniques that can simulate the multiprocessing:
 - ✓ **Process Scheduling:** The OS can rapidly switch between processes, giving the illusion of simultaneous execution. This is known as **context switching**.
 - ✓ **Asynchronous I/O:** When a process waits for I/O (reading from a disk), the operating system can switch to another process, making better use of CPU time.
 - ✓ **Multithreading:** Within a single process, multiple threads can be created to execute different tasks concurrently. This can improve responsiveness and resource utilization, even on a uniprocessor.
 - ✓ **Distributed Computing:** While not strictly multiprocessing, distributed computing involves using multiple computers connected by a network to work on a single task. This can simulate the behavior of a multiprocessor system.

- Computer runs many processes simultaneously

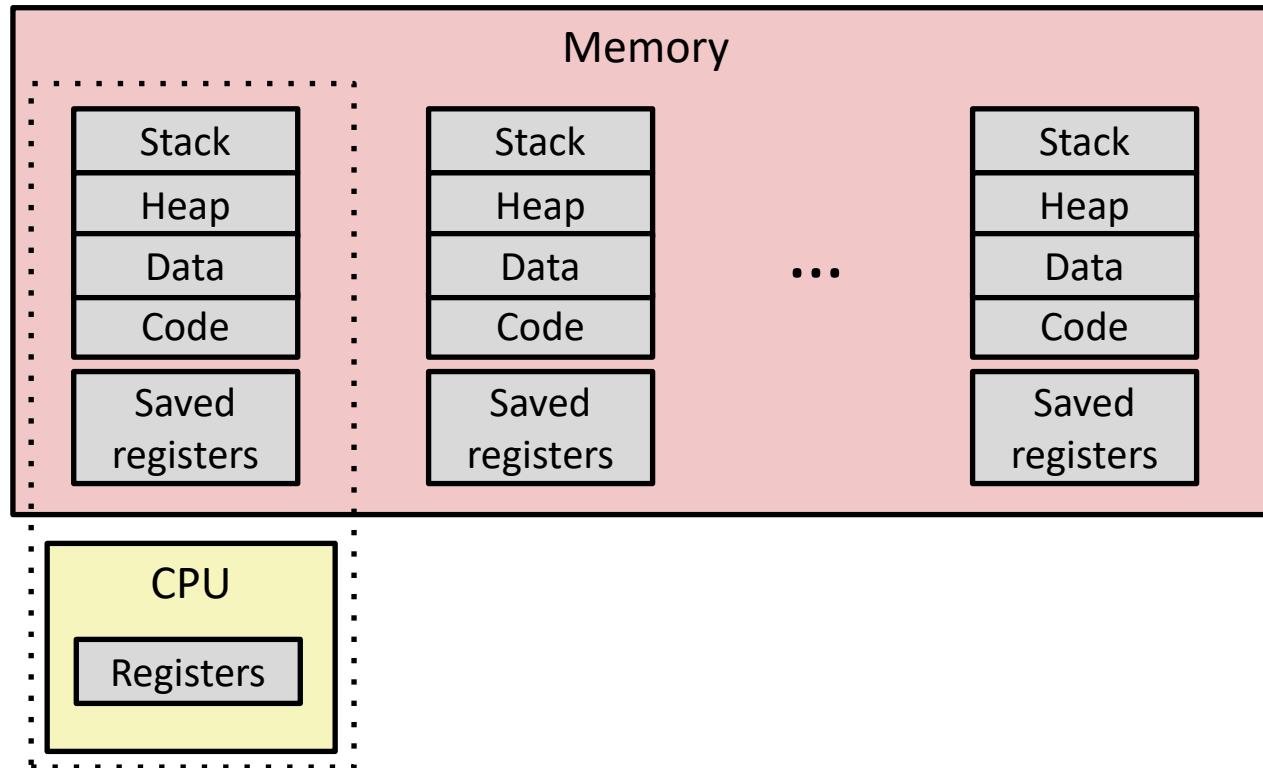
- Applications for one or more users
 - Web browsers, email clients, editors, ...
- Background tasks
 - Monitoring network & I/O devices

} user-level

} mostly kernel/OS - level

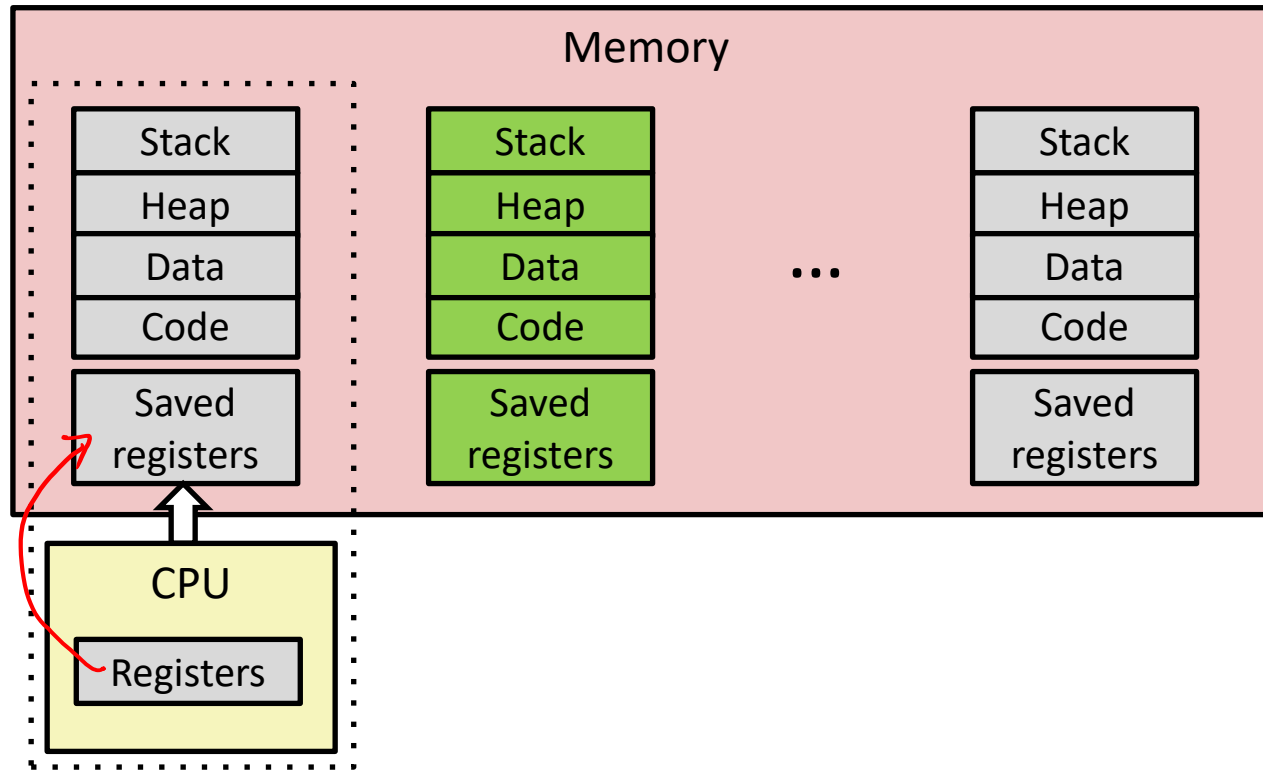


Multiprocessing: The Reality



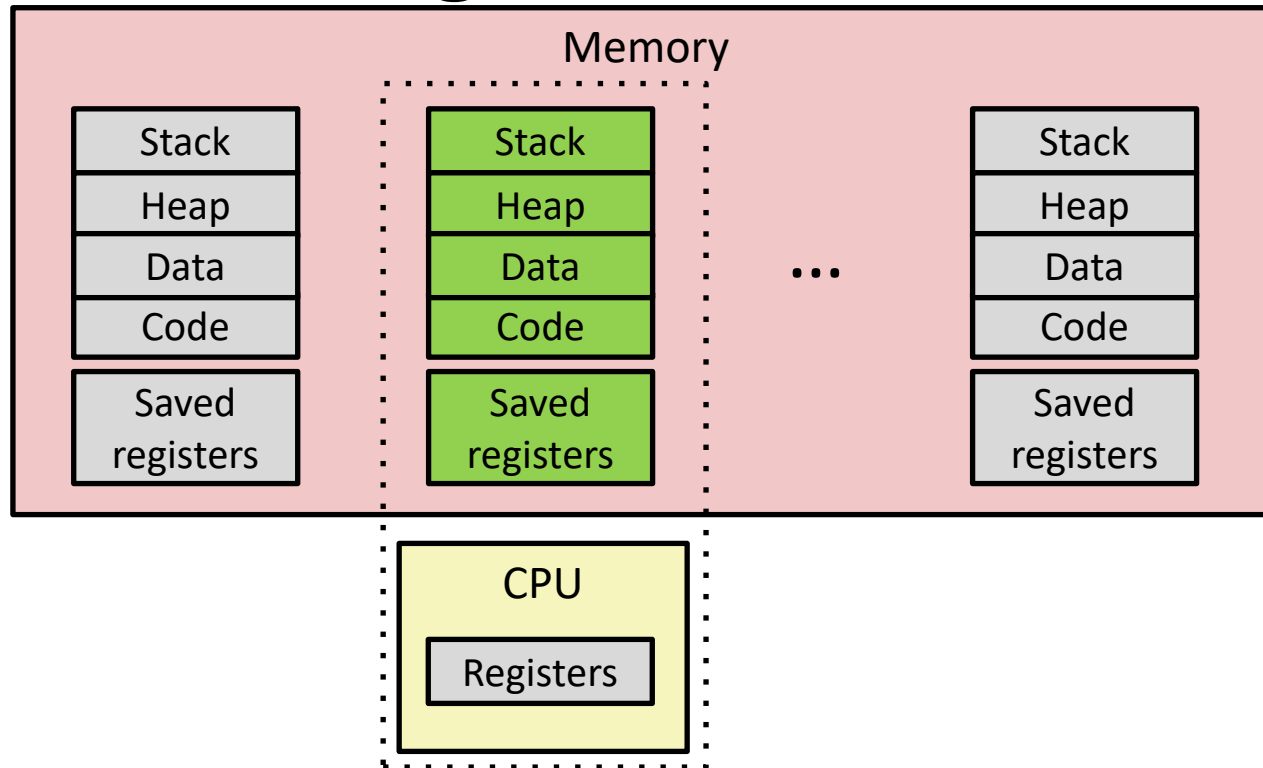
- Single processor executes multiple processes *concurrently*
 - Process executions interleaved, CPU runs *one at a time*
 - Address spaces managed by virtual memory system (later in course)
 - *Execution context* (register values, stack, ...) for other processes saved in memory

Multiprocessing



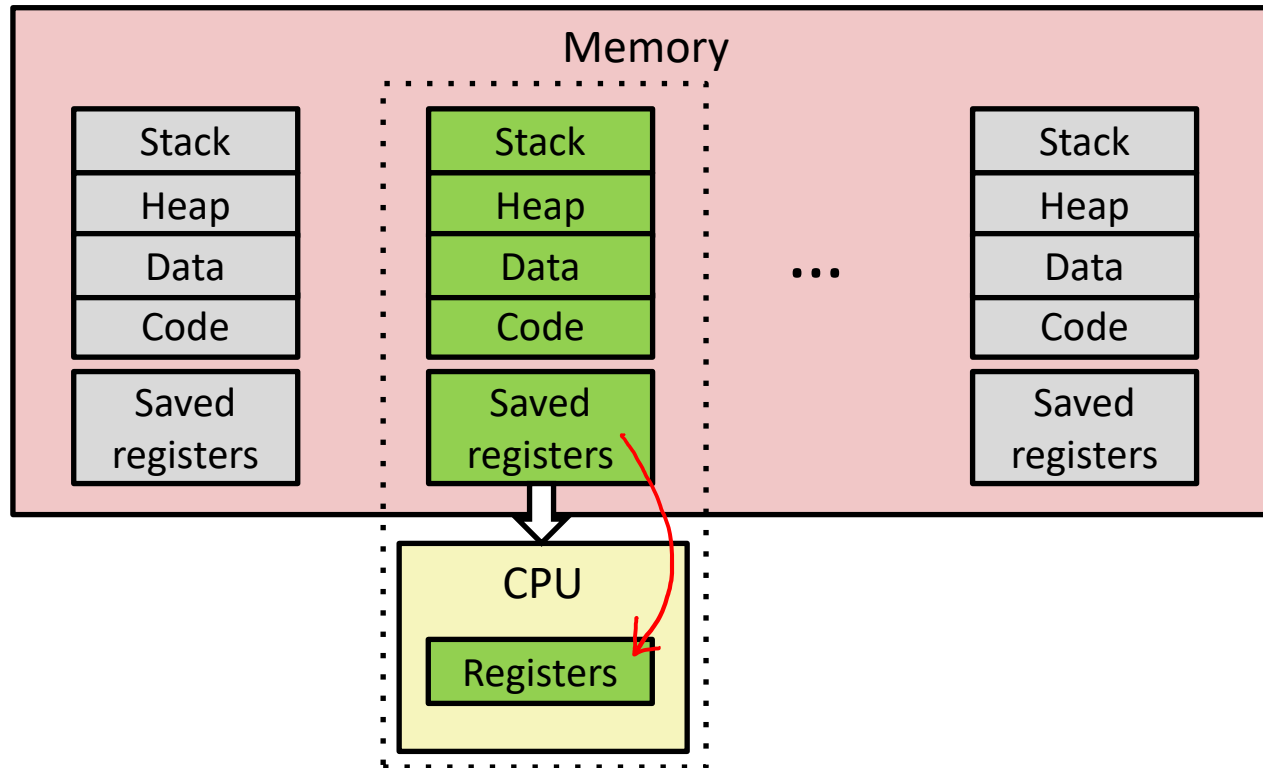
- Context switch
 - 1) Save current registers in memory

Multiprocessing



- Context switch
 - 1) Save current registers in memory
 - 2) Schedule next process for execution (OS decides)

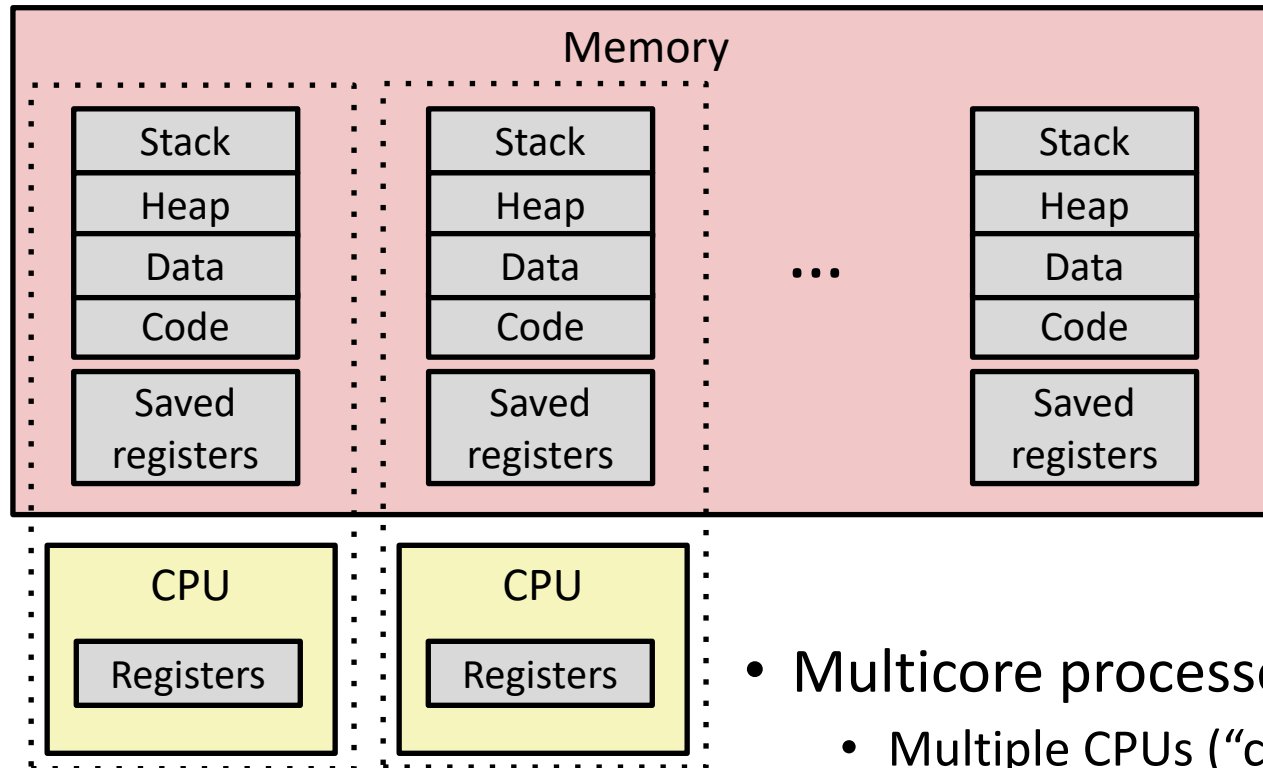
Multiprocessing



❖ Context switch

- 1) Save current registers in memory
- 2) Schedule next process for execution
- 3) **Load saved registers and switch address space**

Multiprocessing: The (Modern) Reality

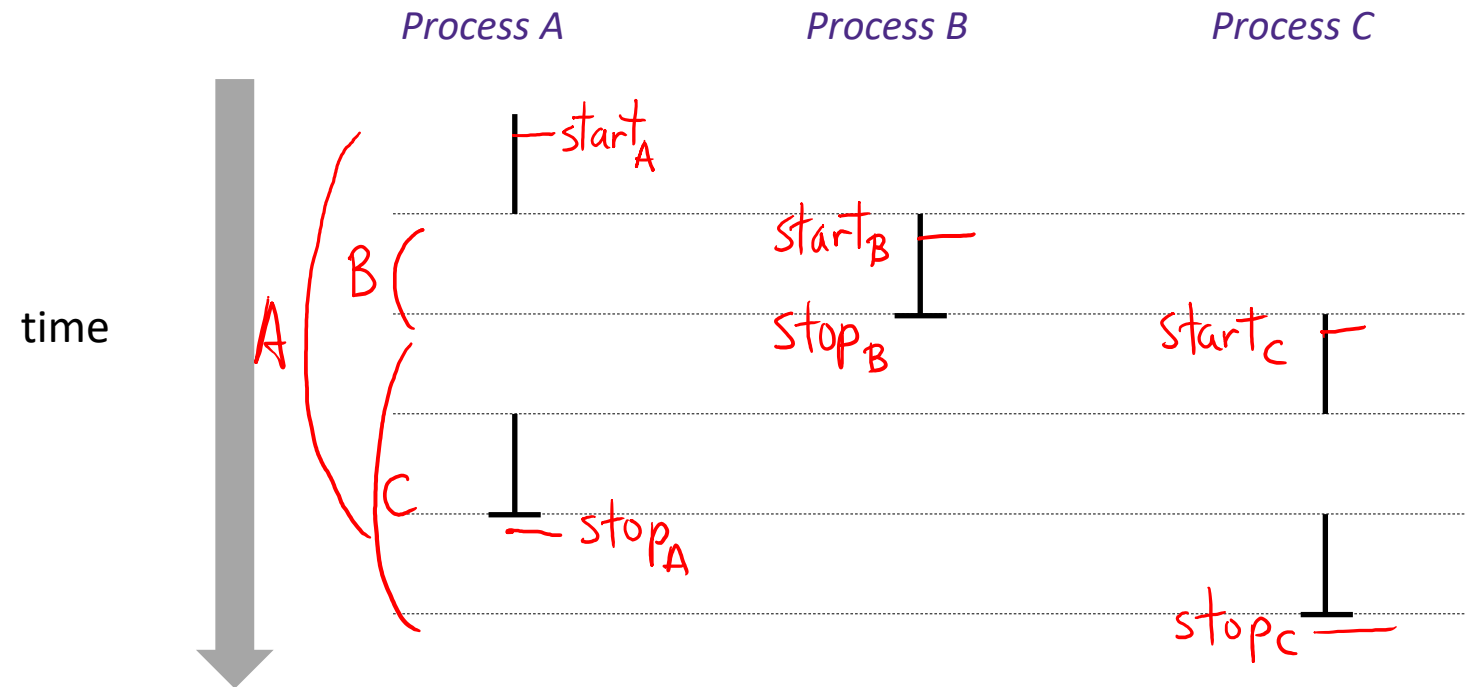


- Multicore processors
 - Multiple CPUs (“cores”) on single chip
 - Share main memory (and some of the caches)
 - Each can execute a separate process
 - Kernel schedules processes to cores
 - **Still constantly swapping processes**

Concurrent Processes

- Each process is a logical control flow
- Two processes *run concurrently* if their instruction executions overlap in time
 - Otherwise, they are *sequential*
- Example: (running on single core)
 - Concurrent: A & B, A & C
 - Sequential: B & C

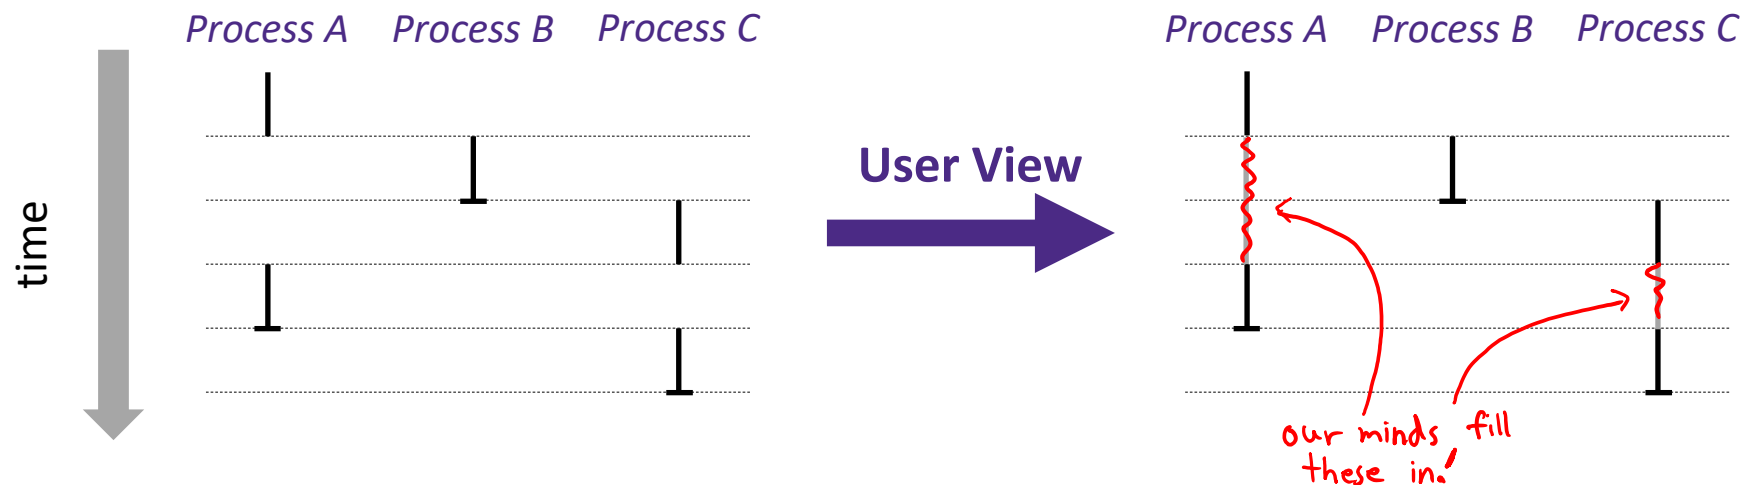
Assume only one CPU



User's View of Concurrency

Assume only one CPU

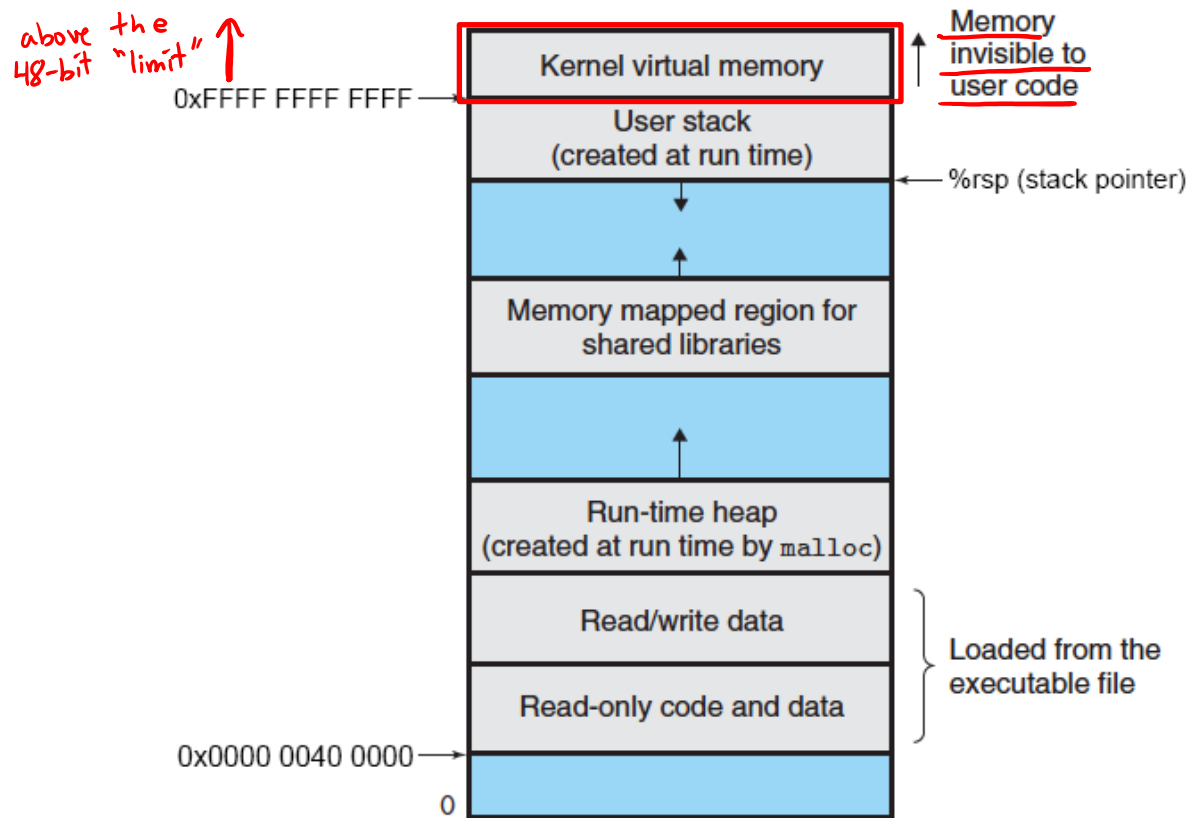
- Control flows for concurrent processes are physically disjoint in time
 - CPU only executes instructions for one process at a time
- However, the user can *think of* concurrent processes as executing at the same time, in *parallel*



Context Switching

Assume only one CPU

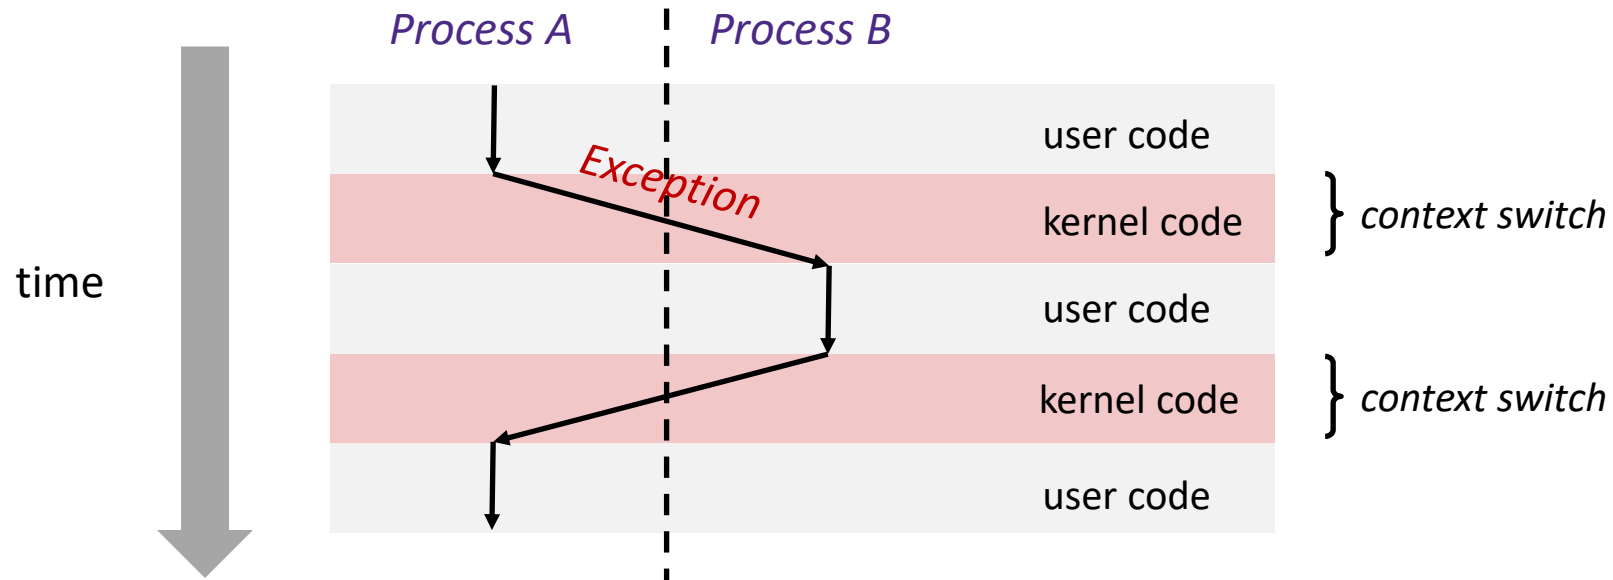
- Processes are managed by a *shared* chunk of OS code called the **kernel**
 - The kernel is not a separate process, but rather runs as part of a user process
- In x86-64 Linux:
 - Same address in each process refers to same shared memory location



Context Switching

Assume only one CPU

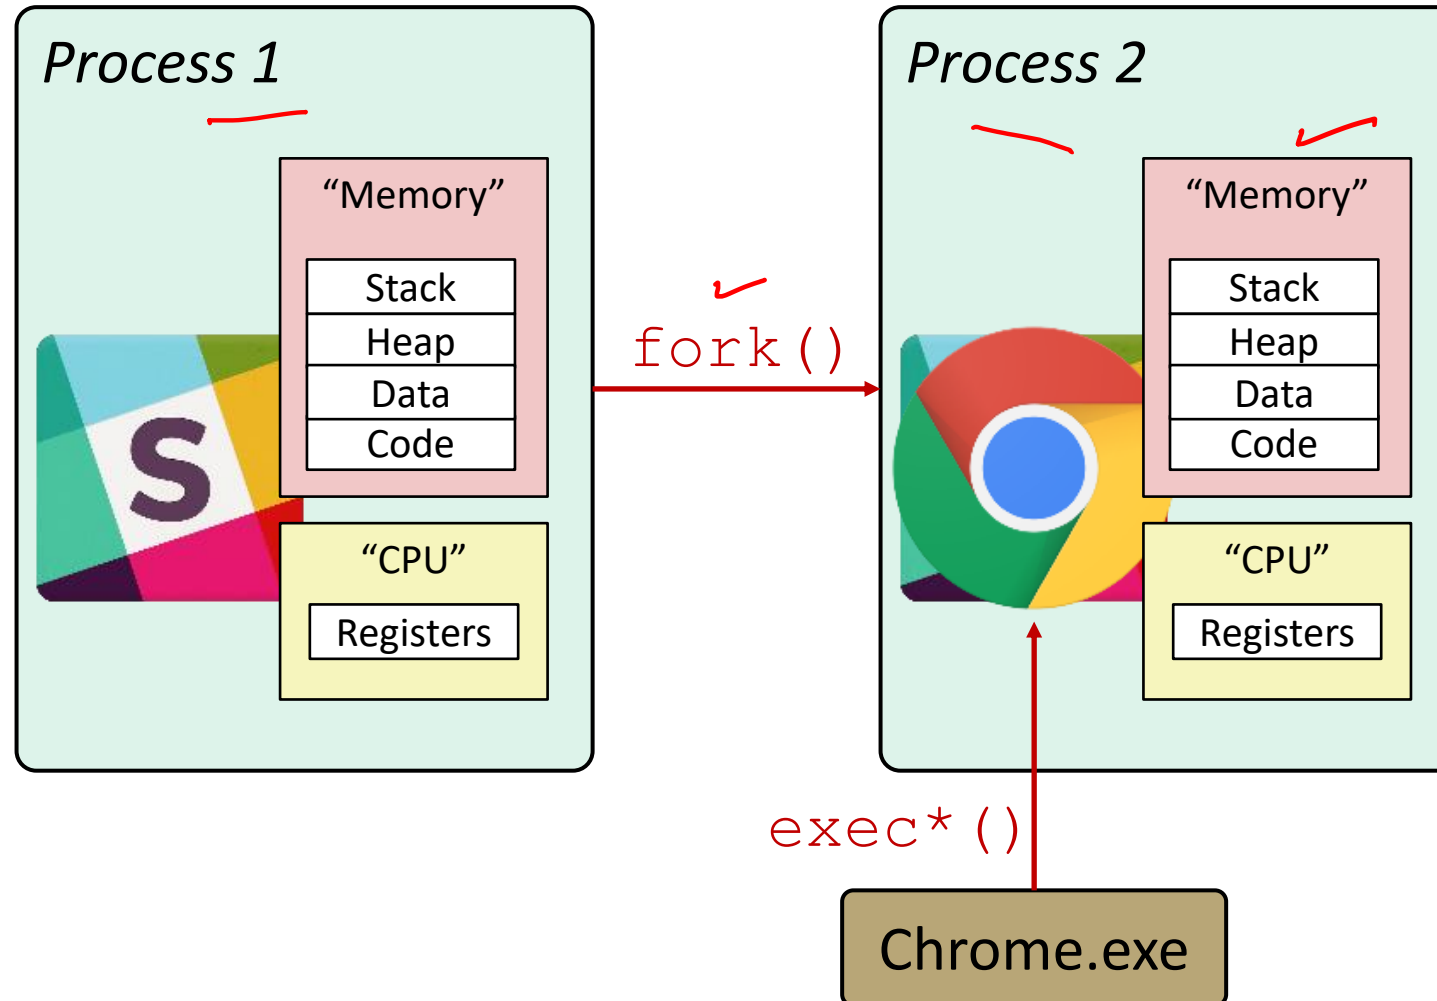
- Processes are managed by a *shared* chunk of OS code called the **kernel**
 - The kernel is not a separate process, but rather runs as part of a user process
- Context switch passes control flow from one process to another and is performed using kernel code



Processes

- Processes and context switching
- **Creating new processes**
 - `fork()` , `exec*()` , and `wait()`
- Zombies

Creating New Processes & Programs



Creating New Processes & Programs

- [✓]fork-[✓]exec model (Linux):
 - fork() creates a copy of the current process
 - exec*() replaces the current process' code and address space with the code for a different program
 - Family: execv, execl, execve, execl, execvp, execlp
 - fork() and execve() are *system calls*
- Other system calls for process management:
 - getpid()
 - exit()
 - wait(), waitpid()


fork: Creating New Processes

```
pid_t pid = fork() ;  
if (pid == 0) {  
    printf("hello from child\n") ;  
} else {  
    printf("hello from parent\n") ;  
}
```

- **pid_t fork(void)**
 - Creates a new “**child**” process that is *identical* to the calling “**parent**” process, including all state (memory, registers, etc.)
 - Returns 0 to the **child** process
 - Returns child’s **process ID (PID)** to the **parent** process
- Child is *almost* identical to parent:
 - Child gets an identical (but separate) copy of the parent’s virtual address space
 - Child has a different PID than the parent
- fork is unique (and often confusing) because it is called **once** but returns “**twice**”

Understanding fork

Process X (parent)




```
pid_t pid = fork();  
if (pid == 0) {  
    printf("hello from child\n");  
} else {  
    printf("hello from parent\n");  
}
```

PID X

fork

Process Y (child)



```
pid_t pid = fork();  
if (pid == 0) {  
    printf("hello from child\n");  
} else {  
    printf("hello from parent\n");  
}
```

PID Y

Understanding fork

Process X (parent)

➔

```
pid_t pid = fork();
if (pid == 0) {
    printf("hello from child\n");
} else {
    printf("hello from parent\n");
}
```

Process Y (child)

➔

```
pid_t pid = fork();
if (pid == 0) {
    printf("hello from child\n");
} else {
    printf("hello from parent\n");
}
```

PID X

➔

```
pid_t pid = fork();
if (pid == 0) {
    printf("hello from child\n");
} else {
    printf("hello from parent\n");
}
```

pid = Y

PID Y

➔

```
pid_t pid = fork();
if (pid == 0) {
    printf("hello from child\n");
} else {
    printf("hello from parent\n");
}
```

pid = 0

Understanding fork

Process X (parent)

➔

```
pid_t pid = fork();
if (pid == 0) {
    printf("hello from child\n");
} else {
    printf("hello from parent\n");
}
```

➔

```
pid_t pid = fork();
if (pid == 0) {
    printf("hello from child\n");
} else {
    printf("hello from parent\n");
}
```

pid = Y

hello from parent

Process Y (child)

➔

```
pid_t pid = fork();
if (pid == 0) {
    printf("hello from child\n");
} else {
    printf("hello from parent\n");
}
```

➔

```
pid_t pid = fork();
if (pid == 0) {
    printf("hello from child\n");
} else {
    printf("hello from parent\n");
}
```

pid = 0

hello from child

Which one appears first?
non-deterministic!

Fork Example

```
void fork1() {  
    int x = 1;  
    pid_t pid = fork();  
    if (pid == 0) splits here  
        printf("Child has x = %d\n", ++x); child only  
    else  
        printf("Parent has x = %d\n", --x); parent only  
    printf("Bye from process %d with x = %d\n", getpid(), x); both  
}
```

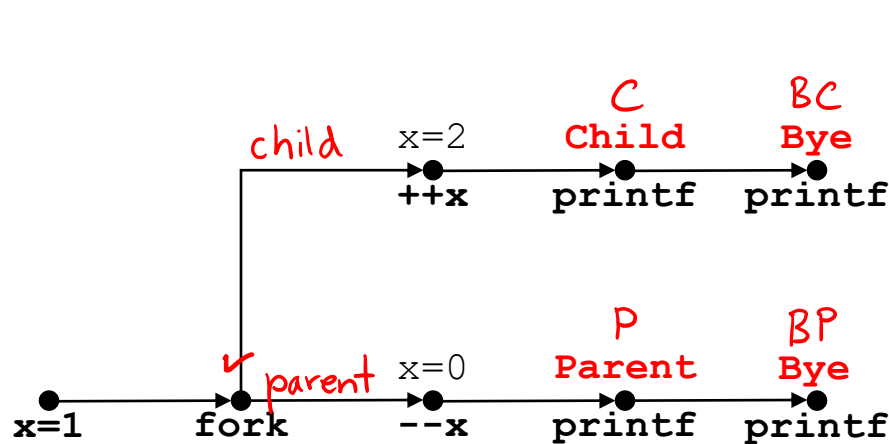
- Both processes continue/start execution after `fork`
 - Child starts at instruction after the call to `fork` (storing into `pid`)
- Can't predict execution order of parent and child
- Both processes start with `x=1`
 - Subsequent changes to `x` are independent
- Shared open files: `stdout` is the same in both parent and child

Modeling fork with Process Graphs

- A *process graph* is a useful tool for capturing the partial ordering of statements in a concurrent program
 - Each vertex is the execution of a statement
 - $a \rightarrow b$ means a happens before b
 - Edges can be labeled with current value of variables
 - **printf** vertices can be labeled with output
 - Each graph begins with a vertex with no in-edges
- Any *topological sort* of the graph corresponds to a feasible total ordering
 - Total ordering of vertices where all edges point from left to right

Fork Example: Possible Output

```
void fork1() {
    int x = 1;
    pid_t pid = fork();
    if (pid == 0)
        printf("Child has x = %d\n", ++x);
    else
        printf("Parent has x = %d\n", --x);
    printf("Bye from process %d with x = %d\n", getpid(), x);
}
```



Possible

C	P	C	C
BC	BP	P	P
P	C	BC	BP
BP	BP	BP	BC

etc...

Not Possible

C	P
BC	BC
BP	C
P	BP

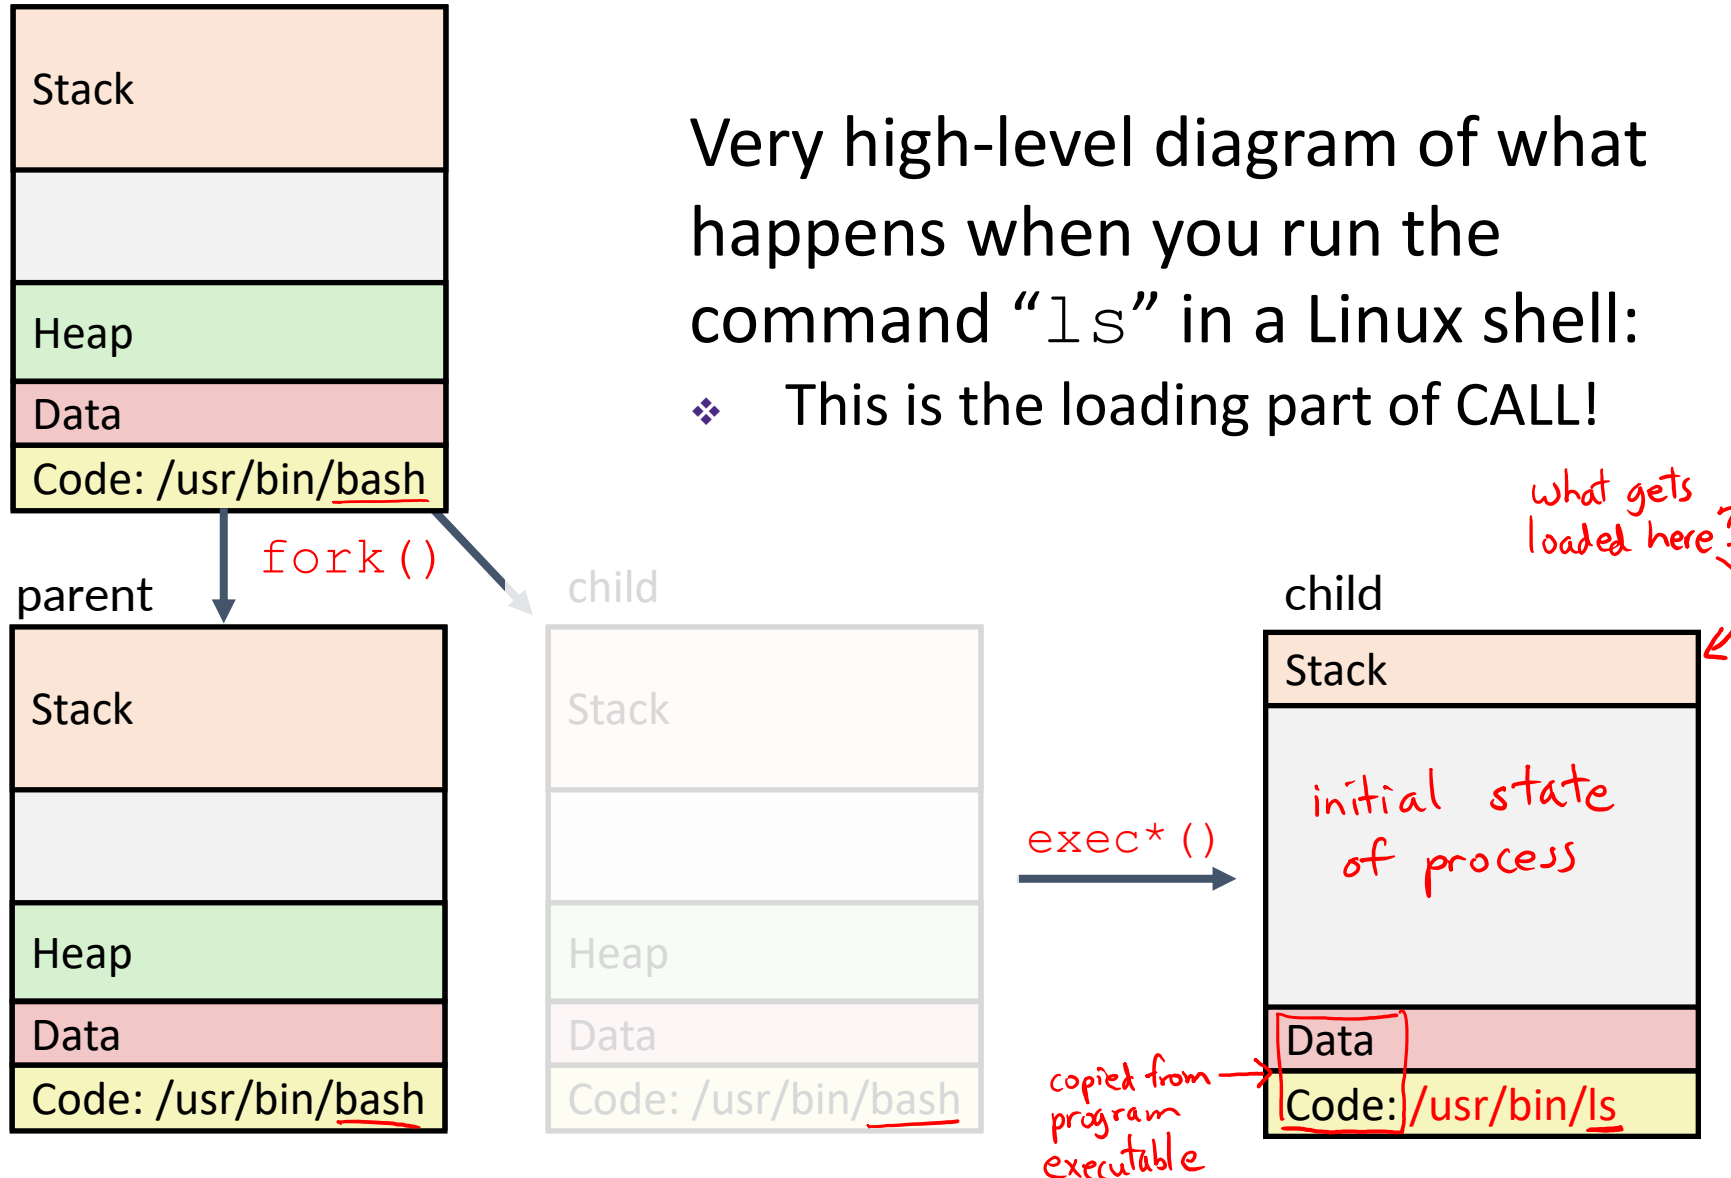
etc...

as long as C comes before BC
and P comes before BP

- fork-exec model:
 - `fork()` creates a copy of the current process
 - `exec*()` replaces the current process' code and address space with the code for a different program
 - Whole family of `exec` calls – see **`exec(3)`** and **`execve(2)`**

```
// Example arguments: path="/usr/bin/ls",  
//      argv[0]="/usr/bin/ls", argv[1]="-ahl", argv[2]=NULL  
void fork_exec(char *path, char *argv[]) {  
    pid_t pid = fork();  
    if (pid != 0) {  
        printf("Parent: created a child %d\n", pid);  
    } else {  
        printf("Child: about to exec a new program\n");  
        execv(path, argv);  
    }  
    printf("This line printed by parent only!\n");  
}
```

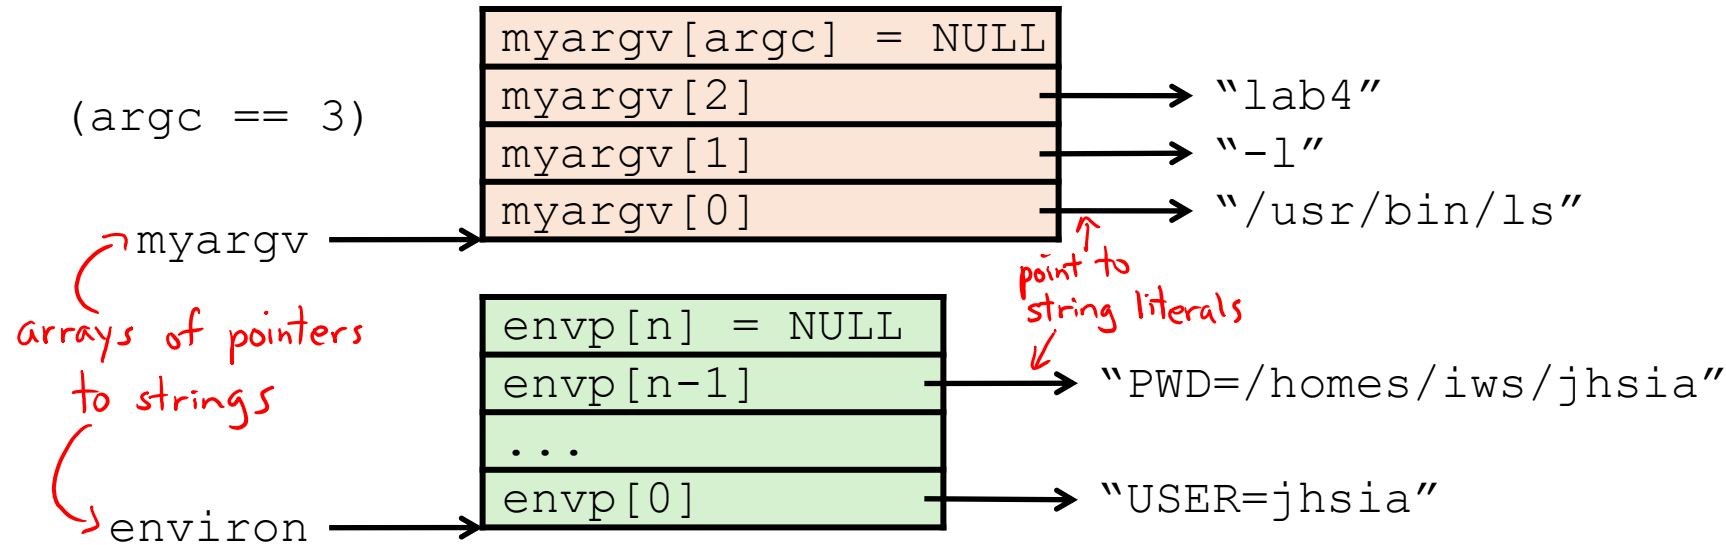
Exec-ing a new program



execve Example

`int main(int argc, char* argv[])`
get command-line arguments into program

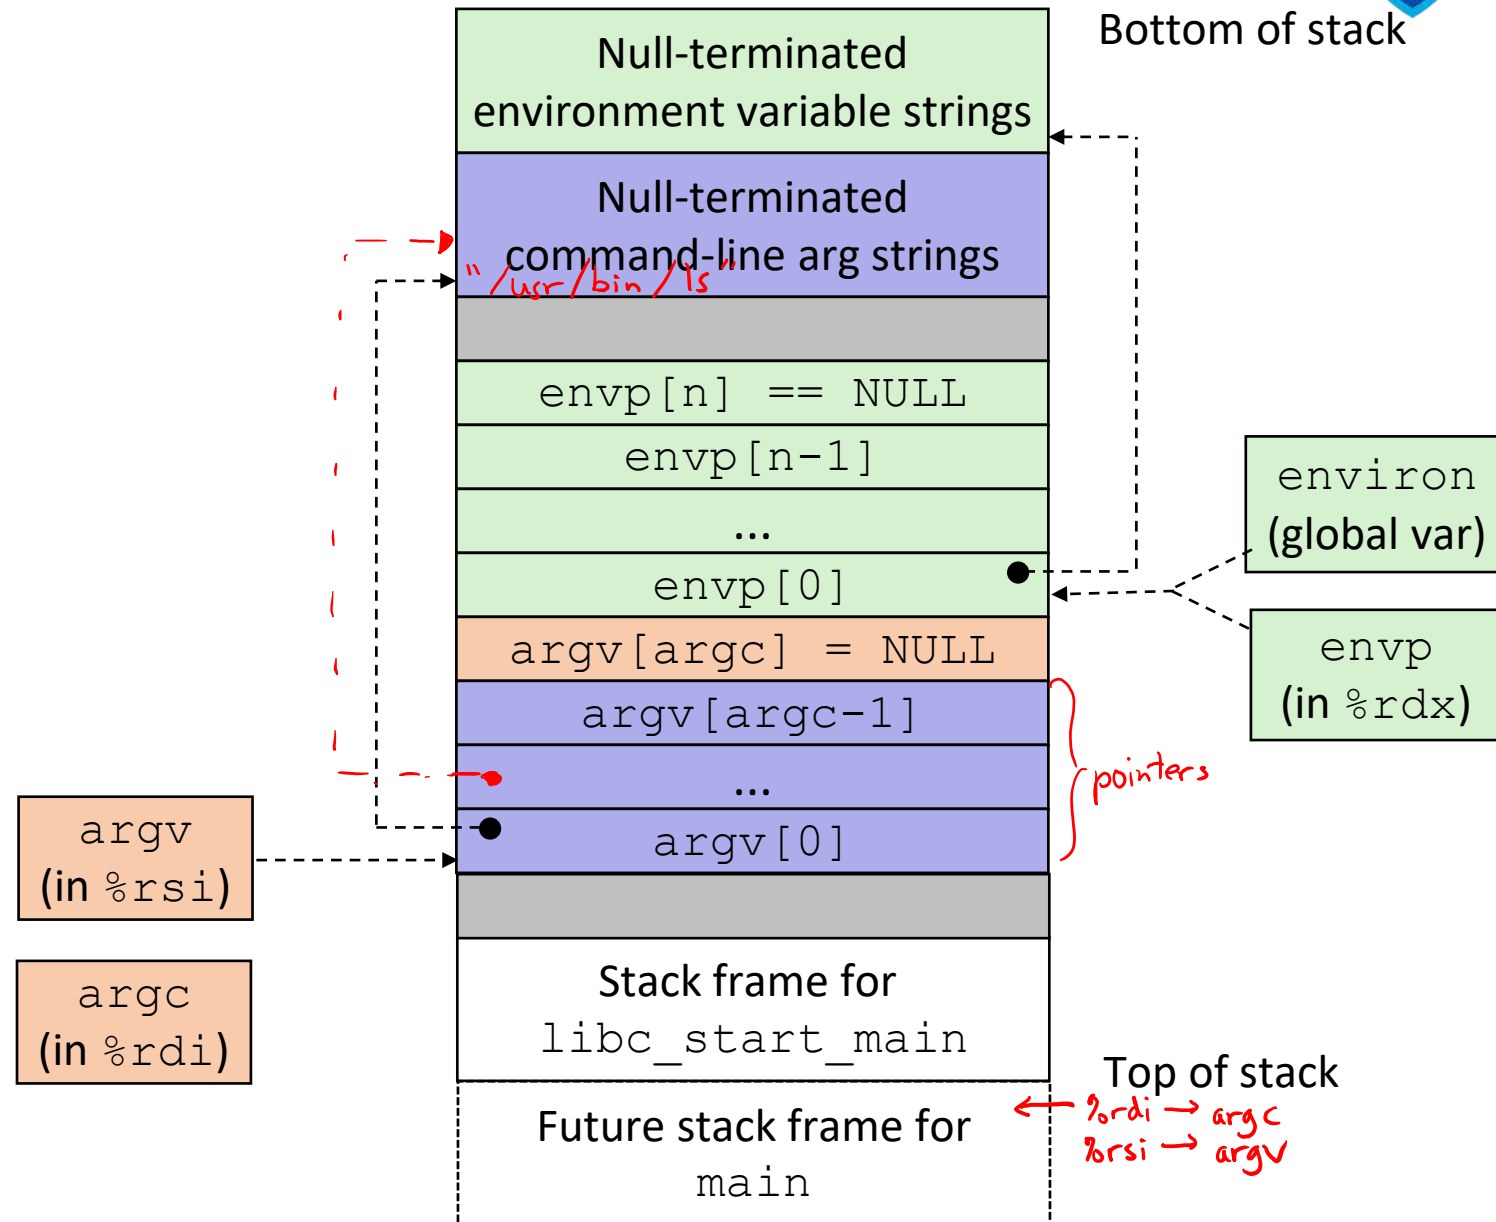
Execute `"/usr/bin/ls -l lab4"` in child process using current environment:



```
if ((pid = fork()) == 0) { /* Child runs program */
    if (execve(myargv[0], myargv, environ) < 0) {
        printf("%s: Command not found.\n", myargv[0]);
        exit(1);
    }
}
```

Run the `printenv` command in a Linux shell to see your own environment variables

Structure of the Stack when a new program starts



exit: Ending a process

- **void** exit(**int** status)
 - Exits a process
 - Status code: 0 is used for a normal exit, nonzero for abnormal exit

Processes

- Processes and context switching
- Creating new processes
 - `fork()`, `exec*()`, and `wait()`
- **Zombies**

Zombies

- When a process terminates, it still consumes system resources
 - Various tables maintained by OS
 - Called a “**zombie**” (a living corpse, half alive and half dead)
- *Reaping* is performed by parent on terminated child
 - Parent is given exit status information and kernel then deletes zombie child process
- What if parent doesn't reap?
 - If any parent terminates without reaping a child, then the orphaned child will be reaped by `init` process (`pid == 1`)
 - **Note:** on more recent Linux systems, `init` has been renamed **`systemd`**
 - In long-running processes (e.g. shells, servers) we need *explicit* reaping

wait: Synchronizing with Children

- **int** wait(**int** *child_status)
 - Suspends current process (*i.e.* the parent) until one of its children terminates
 - Return value is the PID of the child process that terminated
 - *On successful return, the child process is reaped*
 - If child_status != NULL, then the *child_status value indicates why the child process terminated
 - Special macros for interpreting this status – see **man wait(2)**
- **Note:** If parent process has multiple children, wait will return when *any* of the children terminates
 - waitpid can be used to wait on a specific child process

wait: Synchronizing with Children


```
void fork_wait() {
    int child_status;

    if (fork() == 0) {
        printf("HC: hello from child\n");
        exit(0);
    } else {
        printf("HP: hello from parent\n");
        wait(&child_status);
        printf("CT: child has terminated\n");
    }
    printf("Bye\n");
}
```

} child

} parent

forks.c



Feasible output:

```
HC  HP
HP  HC
CT  CT
Bye Bye
```

Infeasible output:

```
HP
CT
Bye
HC
```

Example: Zombie

```
void fork7() {  
    if (fork() == 0) {  
        /* Child */  
        printf("Terminating Child, PID = %d\n",  
               getpid());  
        exit(0);  
    } else {  
        printf("Running Parent, PID = %d\n",  
               getpid());  
        while (1); /* Infinite loop */  
    }  
}
```

parent persists *forks.c*

```
linux> ./forks 7 &  
[1] 6639  
Running Parent, PID = 6639  
Terminating Child, PID = 6640  
linux> ps  
  PID TTY          TIME CMD  
 6585 ttyp9        00:00:00 tcsh  
 6639 ttyp9        00:00:03 forks  
 6640 ttyp9        00:00:00 forks <defunct>  
 6641 ttyp9        00:00:00 ps  
linux> kill 6639  
[1] Terminated  
linux> ps  
  PID TTY          TIME CMD  
 6585 ttyp9        00:00:00 tcsh  
 6642 ttyp9        00:00:00 ps
```

- ps shows child process as "defunct"

- Killing parent allows child to be reaped by init

Example:

Non-terminating Child

```
void fork8() {  
    if (fork() == 0) {  
        /* Child */  
        printf("Running Child, PID = %d\n",  
               getpid());  
        while (1); /* Infinite loop */  
    } else {  
        printf("Terminating Parent, PID = %d\n",  
               getpid());  
        exit(0);  
    }  
}
```

forks.c

```
linux> ./forks 8  
Terminating Parent, PID = 6675  
Running Child, PID = 6676  
linux> ps  
  PID TTY          TIME CMD  
 6585 tttyp9        00:00:00 tcsh  
 6676 tttyp9        00:00:06 forks  
 6677 tttyp9        00:00:00 ps  
linux> kill 6676  
linux> ps  
  PID TTY          TIME CMD  
 6585 tttyp9        00:00:00 tcsh  
 6678 tttyp9        00:00:00 ps
```

- Child process still active even though parent has terminated
- Must kill explicitly, or else will keep running indefinitely

Process Management Summary

- `fork` makes two copies of the same process (parent & child)
 - Returns different values to the two processes
- `exec*` replaces current process from file (new program)
 - Two-process program:
 - First `fork()`
 - `if (pid == 0) { /* child code */ } else { /* parent code */ }`
 - Two different programs:
 - First `fork()`
 - `if (pid == 0) { execv(...) } else { /* parent code */ }`
- `wait` or `waitpid` used to synchronize parent/child execution and to reap child process

Summary

- Processes
 - At any given time, system has multiple active processes
 - On a one-CPU system, only one can execute at a time, but each process appears to have total control of the processor
 - OS periodically “context switches” between active processes
 - Implemented using *exceptional control flow*
- Process management
 - `fork`: one call, two returns
 - `execve`: one call, usually no return
 - `wait` or `waitpid`: synchronization
 - `exit`: one call, no return



Thanks

Q & A