

Processes in Linux

What is a Process?

- A process is a **program in execution**.
- It is a **dynamic entity** managed by the operating system.

Components of a Process

A process consists of:

- **Program Counter (PC):** Tracks the address of the next instruction.
 - **CPU Registers:** Includes stack pointer, general-purpose registers, etc.
 - **Process Stack:** Stores function parameters, local variables, and return addresses.
 - **Memory Regions:** Text, stack, and heap.
-

Process Memory Layout

Every process in Linux has 3 main memory regions:

1. **Text Region:**
 - Contains the **program code**.
 - **Static** in size.
2. **Stack Region:**
 - Stores local variables, function parameters, and return addresses.
 - Grows **downward**.
3. **Data Region (Heap):**
 - Stores **program data** that grows dynamically as per user requests.
 - Grows **upward**.

Diagram:

mathematica

```
0x0000
|   Text Region
|   Data (Heap)
|   Unused Free Space
|   Stack Region
0xffff
```

User Stack vs Kernel Stack

1. **User Stack:**
 - Resides in **user space**.
 - Manipulated by the **process itself**.

2. Kernel Stack:

- Resides in **kernel space**.
 - Used by the OS to handle **system calls** and **interrupts**.
 - Manipulated by the **OS**.
-

Process Descriptor (**task_struct**)

- A **data structure** that holds all information about a process.
 - Contains:
 - **State**: Current state of the process.
 - **PID (Process ID)**: Unique ID for each process.
 - **Program Counter**: Tracks the execution point.
 - **Registers**: CPU state information.
 - **Memory Regions**: Text, stack, heap.
 - **Open Files**: List of open files for the process.
 - **Pointers**: To manage process queues.
-

Foreground vs Background Processes

1. Foreground Process:

- Requires **user input** to run.
- Example: Text editors like `vim`.

2. Background Process:

- Runs **continuously** without requiring user input.
- Example: Download commands like `wget`.

Move Process to Background:

- Add `&` at the end of the command:

```
bash
```

```
wget http://example.com/file.iso &
```

Bring Back to Foreground:

- Use `fg` command:

```
bash
```

```
jobs          # Check background jobs  
fg [job_number]
```

Process States

1. **New:** Process is created.
 2. **Ready:** Process is ready to execute.
 3. **Running:** Process is currently being executed.
 4. **Waiting:** Process is waiting for an event/resource.
 5. **Terminated (Zombie):** Process is halted, but its descriptor exists.
-

Process Management Commands

1. View Processes:

- `ps -ef`: Displays all processes.
- `top`: Dynamic display of processes.

2. Kill a Process:

- **Find PID:** Use `ps` or `top`.
- **Kill Command:**

```
bash
```

```
kill [PID]          # Graceful termination
kill -9 [PID]       # Force termination (SIGKILL)
```

3. Kill Multiple Processes:

- Use `killall` to terminate multiple processes:

```
bash
```

```
killall [process_name]
```

- Use `pkill` for specific users:

```
bash
```

```
pkill -u [username]
```

Process Queues

- Processes are organized in **queues**:
 - **Ready Queue:** Jobs ready to execute.
 - **Waiting Queue:** Jobs waiting for an event (I/O, signals).

Example:

```
less
```

```
Ready Queue: PCB A → PCB B
Printer Queue: PCB X → PCB Y
```

Process State Transitions

- A process changes states based on events:
 - **Running** → **Waiting**: Process is waiting for I/O or resources.
 - **Ready** → **Running**: Scheduler selects the process to execute.
 - **Zombie**: Process ends, but its parent has not collected the exit code.
-

Process Creation: Forking

- **Forking** creates a new process in Linux.
- The parent process spawns a child process.

Example Code:

```
c

#include <stdio.h>
#include <unistd.h>

void main() {
    int pid;
    pid = fork(); // Creates child process
    if (pid == 0) {
        printf("Child process\n");
    } else {
        printf("Parent process\n");
    }
}
```

Process Termination

Processes terminate in multiple ways:

1. **exit()**: The process calls exit explicitly.
 2. **kill()**: Parent process or user sends a kill signal.
 3. **SIGKILL**: Force termination with `kill -9`.
-

Zombie Processes

- **Definition**: A process that has completed execution but still has an entry in the process table.
 - **Why?**: Parent process hasn't collected the child's exit status using `wait()`.
 - **Solution**: The `init` process cleans up zombie processes.
-

init Process

- The **first process** started by the OS.

- It **adopts orphan processes** and cleans up zombie processes.
-

Process Representation: `task_struct`

In Linux, a process is represented by the **`task_struct`** structure. It contains:

- **State:** Current status of the process.
 - **Scheduling Information:** Priority and scheduling queues.
 - **Identifiers:** PID, user ID, group ID.
 - **IPC:** Signals, pipes, semaphores.
 - **File System:** Open files, root, and working directories.
 - **Timers:** Process-specific timers.
 - **Virtual Memory:** Memory mappings.
-

Important Commands Summary

Command	Description
<code>ps</code>	Displays processes running on the system.
<code>top</code>	Shows dynamic list of processes.
<code>kill [PID]</code>	Terminates a specific process.
<code>killall [name]</code>	Kills all processes by name.
<code>fg</code>	Brings a background process to the foreground.
<code>jobs</code>	Lists background jobs.
<code>pstree</code>	Displays process hierarchy.

Fork and Zombie Processes in Linux

What is a Process?

- A **process** is a program in **execution**.
 - It consists of:
 - **Registers**: CPU state information.
 - **Memory**: Code, data, stack, and heap regions.
 - **Process Context**: Includes open files, stack pointers, etc.
-

Process Management

1. **Multiprocessing**: Simultaneous execution of multiple processes.
 - On a **single CPU**: The OS rapidly switches between processes using **context switching**.
 - On a **multicore CPU**: Multiple cores can execute different processes in parallel.
 2. **Context Switching**:
 - The OS saves the current process state (registers, program counter, etc.) and loads the next process.
 - Provides the **illusion** of parallel execution.
-

Creating New Processes

1. fork() System Call

- **Definition**: `fork()` creates a **child process** by duplicating the **parent process**.
- The child is an **exact copy** of the parent but has a **different Process ID (PID)**.
- Returns:
 - `0` to the **child process**.
 - Child's PID to the **parent process**.

Example Code:

```
c

#include <stdio.h>
#include <unistd.h>

int main() {
    pid_t pid = fork();

    if (pid == 0) {
        printf("Hello from Child\n");
    } else {
        printf("Hello from Parent\n");
    }
}
```

```
    }  
    return 0;  
}
```

Output

- "Hello from Parent" and "Hello from Child" can appear in **any order** due to concurrent execution.
-

2. exec() System Call

- **Definition:** `exec()` replaces the **current process's code** and address space with a **new program**.
- `fork()` → creates a process, then `exec()` → loads a new program.

Types of `exec()`:

- `execv`, `execl`, `execve`, `execle`, `execvp`, `execlp`.

Example:

```
c  
  
#include <stdio.h>  
#include <unistd.h>  
  
int main() {  
    char *args[] = {"/bin/ls", "-l", NULL};  
    execv("/bin/ls", args); // Replaces current process with ls command  
    return 0;  
}
```

3. wait() System Call

- **Definition:** `wait()` makes the parent process **wait** until one of its child processes terminates.
- **Why?:** To **reap** terminated child processes and prevent **zombies**.

Example:

```
c  
  
#include <stdio.h>  
#include <unistd.h>  
#include <sys/wait.h>  
  
int main() {  
    pid_t pid = fork();  
  
    if (pid == 0) {  
        printf("Child Process\n");  
        return 0;  
    } else {  
        wait(NULL);  
    }
```

```
        printf("Parent: Child has terminated\n");
    }
    return 0;
}
```

Zombie Processes

What is a Zombie Process?

- A **zombie process** is a process that has **terminated**, but its **process descriptor** (entry in the process table) is still present.
- **Why?:** The parent process hasn't yet **collected the child's exit status** using `wait()` or `waitpid()`.

Why Do Zombies Exist?

- **Exit status** must be passed to the parent process.
 - The OS keeps the process descriptor until the parent **reaps** it.
-

How to Detect Zombies

- Use the `ps` command to view processes. Zombie processes appear as **<defunct>**.

Example:

```
bash
```

```
ps -ef
```

Output:

```
less
```

```
PID   TTY    TIME CMD
1234  pts/1  0:00 a.out <defunct>
```

How to Handle Zombies

1. Reap Zombies Using `wait()`:

- The parent process should call `wait()` to clean up the terminated child.

2. Orphaned Zombies:

- If a parent process terminates without reaping its child, the **init** process (PID = 1) adopts the child and cleans it up.

3. Kill the Parent Process:

- Use `kill` to terminate the parent. The zombie will then be reaped by `init`.

Example to Create a Zombie:

```
c

#include <stdio.h>
#include <unistd.h>
#include <stdlib.h>

int main() {
    if (fork() == 0) {
        printf("Child Process: PID = %d\n", getpid());
        exit(0); // Child terminates
    } else {
        printf("Parent Process: PID = %d\n", getpid());
        while (1); // Parent process is stuck in an infinite loop
    }
    return 0;
}
```

Output in ps:

The child will appear as **<defunct>** until the parent process terminates or reaps it.

Process Management Summary

1. **fork()**: Creates a new child process.
 2. **exec()**: Replaces the current process with a new program.
 3. **wait()**: Makes the parent wait for the child to terminate.
 4. **Zombie**: Process that has terminated but whose entry remains in the process table.
-

Key Commands

Command	Description
ps	Displays all running processes.
top	Monitors processes dynamically.
kill [PID]	Terminates a process with the given PID.
wait	Makes the parent wait for child termination.
pstree	Displays the process hierarchy.
